

Hbase lab

Invoke Hbase

We have to start daemons and zookeeper then invoke HBase shell

```
(base) [bigdata@localhost ~]$ start-all.sh
```

```
(base) [bigdata@localhost hbase]$ start-hbase.sh
```

```
(base) [bigdata@localhost hbase]$ hbase shell
```

For Reference, please visit: <http://hbase.apache.org/2.0/book.html#shell>

Version 2.3.3, r3e4bf4bee3a08b25591b9c22fea0518686a7e834, Wed Oct 28 06:36:25 UTC 2020

Took 0.0009 seconds

```
hbase(main):001:0> version
```

2.3.3, r3e4bf4bee3a08b25591b9c22fea0518686a7e834, Wed Oct 28 06:36:25 UTC 2020

Took 0.0006 seconds

```
hbase(main):002:0> table_help
```

Help for table-reference commands.

You can either create a table via 'create' and then manipulate the table via commands like 'put', 'get', etc.

However, as of 0.96, you can also get a reference to a table, on which you can invoke commands.

For instance, you can get create a table and keep around a reference to it via:

```
hbase> t = create 't', 'cf'
```

Or, if you have already created the table, you can get a reference to it: `hbase> t = get_table 't'`

You can do things like call 'put' on the table:

```
hbase> t.put 'r', 'cf:q', 'v'
```

which puts a row 'r' with column family 'cf', qualifier 'q' and value 'v' into table t.

To read the data out, you can scan the table:

```
hbase> t.scan
```

which will read all the rows in table 't'.

Essentially, any command that takes a table name can also be done via table reference.

Other commands include things like: get, delete, deleteall,

get_all_columns, get_counter, count, incr. These functions, along with

the standard JRuby object methods are also available via tab completion.

For more information on how to use each of these commands, you can also just type:

```
hbase> t.help 'scan'
```

which will output more information on how to use that command.

You can also do general admin actions directly on a table; things like enable, disable, flush and drop just by typing:

```
hbase> t.enable
```

```
hbase> t.flush
```

```
hbase> t.disable
```

```
hbase> t.drop
```

Note that after dropping a table, your reference to it becomes useless and further usage is undefined (and not recommended).

Took 0.0003 seconds

```
hbase(main):003:0> whoami
```

bigdata (auth:SIMPLE)

groups: bigdata

Took 0.0599 seconds

```
hbase(main):004:0> list
```

TABLE

0 row(s)

```
hbase(main):005:0> status
```

1 active master, 0 backup masters, 1 servers, 0 dead, 2.0000 average load

Took 0.1548 seconds

```
hbase(main):006:0> create 'emp', 'personal data', 'professional data'
```

Created table emp

```
Took 0.7601 seconds
=> Hbase::Table – emp
hbase(main):007:0> list
TABLE
emp
1 row(s)
Took 0.0127 seconds
=> ["emp"]
hbase(main):005:0> create 'course', 'bigdata'
Created table course
Took 0.7574 seconds
=> Hbase::Table – course
hbase(main):006:0> list
TABLE
course
1 row(s)
Took 0.0121 seconds
=> ["course"]
hbase(main):007:0> describe 'course'
Table course is ENABLED
course
COLUMN FAMILIES DESCRIPTION
{NAME => 'bigdata', BLOOMFILTER => 'ROW', IN_MEMORY => 'false', VERSIONS => '1',
KEEP_DELETED_CELLS => 'FALSE', DATA_BLOCK_ENCODING => 'NONE', COMPRESSION
=> 'N
ONE', TTL => 'FOREVER', MIN_VERSIONS => '0', BLOCKCACHE => 'true', BLOCKSIZE =>
'65536', REPLICATION_SCOPE => '0'}
1 row(s)
QUOTAS
0 row(s)
Took 0.2591 seconds
hbase(main):008:0> disable 'course'
Took 0.5467 seconds
hbase(main):009:0> enable 'course'
Took 0.7174 seconds
hbase(main):010:0> show_filters
DependentColumnFilter
KeyOnlyFilter
ColumnCountGetFilter
SingleColumnValueFilter
PrefixFilter
SingleColumnValueExcludeFilter
FirstKeyOnlyFilter
ColumnRangeFilter
ColumnValueFilter
TimestampsFilter
FamilyFilter
QualifierFilter
ColumnPrefixFilter
RowFilter
MultipleColumnPrefixFilter
```

```

InclusiveStopFilter
PageFilter
ValueFilter
ColumnPaginationFilter
Took 0.0084 seconds
hbase(main):011:0> drop 'course'
ERROR: Table course is enabled. Disable it first.
For usage try 'help "drop"'
Took 0.0521 seconds
hbase(main):012:0> is_enabled 'course'
true
Took 0.0082 seconds
=> true
hbase(main):013:0> alter 'course', NAME='learning'
Updating all regions with the new schema...
1/1 regions updated.
Done.
Took 1.7402 seconds
hbase(main):014:0> alter_status 'course'
1/1 regions updated.
Done.
Took 1.0125 seconds
hbase(main):015:0> count 'course', CACHE=>1000
0 row(s)
Took 0.1021 seconds
=> 0
count can fetch 1000 rows at a time from "course" table if it was created and exist.
We can make cache to some lower value if the table consists of more rows. But by default it will fetch
one row at a time.
hbase(main):016:0> count 'course', INTERVAL => 100000
0 row(s)
Took 0.0070 seconds
=> 0

```

Given below is a sample schema of a table named emp. It has two column families: “personal data” and “professional data”.

```

hbase(main):035:0> create 'emp', 'personal data', 'professional data'
Created table emp
Took 1.1530 seconds
=> Hbase::Table - emp
hbase(main):036:0> list
TABLE
course
emp
2 row(s)
Took 0.0150 seconds
=> ["course", "emp"]
hbase(main): 037:0> put 'emp', 1, 'personal data:name', 'Aya'
hbase(main): 038:0> put 'emp', 2, 'personal data:name', 'Marwan'
hbase(main): 039:0> put 'emp', 3, 'personal data:name', 'Suzy'

```

```
hbase(main): 040:0> put 'emp', 4, 'personal data:name' , 'Fahd'
hbase(main): 041:0> put 'emp', 5, 'personal data:name' , 'Ayman'
hbase(main): 042:0> put 'emp', 1, 'personal data:city' , 'Cairo'
hbase(main): 043:0> put 'emp', 2, 'personal data:city' , 'Alex'
hbase(main): 044:0> put 'emp', 3, 'personal data:city' , 'Giza'
hbase(main): 045:0> put 'emp', 4, 'personal data:city' , 'Giza'
hbase(main): 046:0> put 'emp', 5, 'personal data:city' , 'Aswan'
```

```
hbase(main):047:0> scan 'emp'
ROW COLUMN+CELL
0 row(s)
Took 0.0246 seconds
```

```
hbase(main):048:0> put 'emp', 1, 'personal data:age' , '35'
hbase(main):049:0> put 'emp', 2, 'personal data:age' , '33'
hbase(main):050:0> put 'emp', 3, 'personal data:age' , '23'
hbase(main):051:0> put 'emp', 4, 'personal data:age' , '27'
hbase(main):052:0> put 'emp', 5, 'personal data:age' , '38'
```

```
hbase(main):053:0> scan 'emp'
```

```
hbase(main):054:0> put 'emp', 1, 'professional data:position' , 'manager'
hbase(main):055:0> put 'emp', 2, 'professional data:position' , 'team leader'
hbase(main):056:0> scan 'emp', {COLUMNS => 'personal data:name', LIMIT => 2, STARTROW => '3'}
hbase(main):057:0> get 'emp', '1'
```

```
hbase(main):019:0> get 'emp', '1'
COLUMN      CELL
personal data:age  timestamp=2023-03-13T20:39:56.396, value=35
personal data:city  timestamp=2023-03-13T20:37:42.551, value=Cairo
personal data:name  timestamp=2023-03-13T20:32:06.515, value=Aya
professional data:po timestamp=2023-03-13T20:41:11.682, value=manager
sition
1 row(s)
Took 0.0448 seconds
```

```
hbase(main):058:0> get 'emp', '1', {COLUMN => 'personal data:name'}
```

```
hbase(main):020:0> get 'emp', '1', {COLUMN => 'personal data:name'}
COLUMN      CELL
personal data:name  timestamp=2023-03-13T20:32:06.515, value=Aya
1 row(s)
Took 0.0481 seconds
```

```
base(main):060:0> scan 'emp', { COLUMNS => 'personal data:name' , FILTER =>
"ValueFilter(=, 'binary:Aya')" }
ROW          COLUMN+CELL
1           column=personal data:name, timestamp=2022-02-03T21:41:49.484, value=heba
1 row(s)
Took 0.0429 seconds
```

```
hbase(main):021:0> scan 'emp', { COLUMNS => 'personal data:name' ,
FILTER => "ValueFilter(=, 'binary:Aya')" }
ROW          COLUMN+CELL
1           column=personal data:name, timestamp=2023-03-13T20:32:06.5
15, value=Aya
1 row(s)
Took 0.0475 seconds
```

Row key	personal data			professional data	
	personal data			professional data	
Rowkey	Name	City	age	Rowkey	position
1	Aya	Cairo	35	1	manager
2	Marwan	Alex	33	2	Team leader
3	Suzy	Giza	23		
4	Fahd	Giza	27		
5	Ayman	Aswan	38		

hbase(main):038:0> exists 'student'

Table student does not exist

Took 0.0102 seconds

=> false

hbase(main):039:0> exists 'emp'

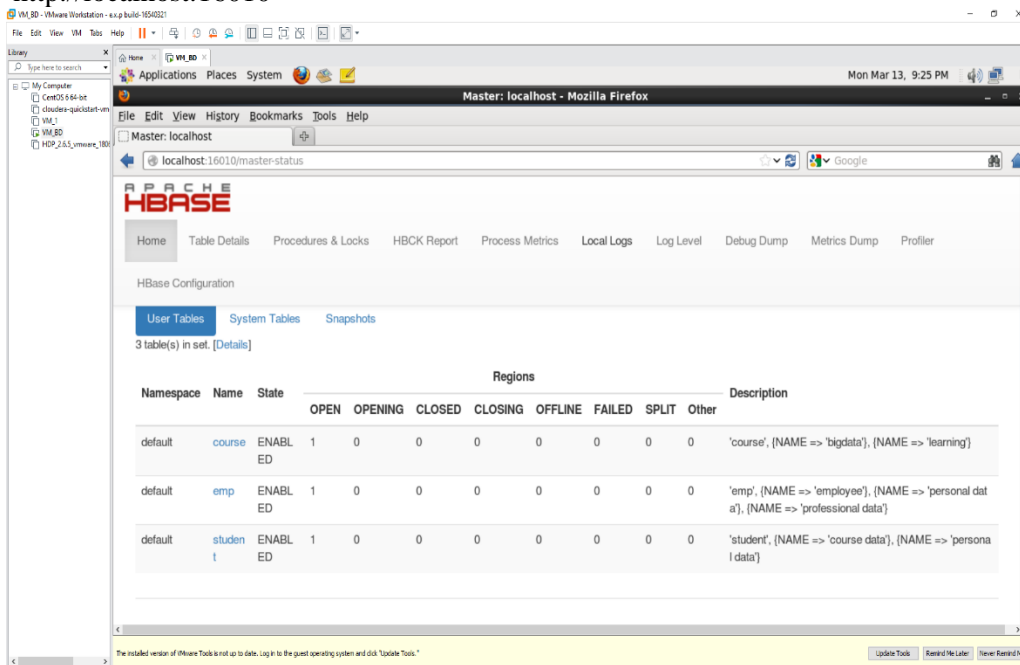
Table emp does exist

Took 0.0114 seconds

=> true

Start Hbase web interface :

<http://localhost:16010>



hbase(main):040:0> exit

hbase\$./bin/stop-hbase.sh

bigdata@localhost\$ stop-all.sh

--The HBase directory can be seen through HDFS as shown

--By exploring HBase we can view and monitor its content with details as permissions owner, group size , date and replication if exist.

Tables Managements commands

These commands will allow programmers to create tables and table schemas with rows and column families. The following are Table Management commands

- Create
- List
- Describe
- Disable
- Disable_all
- Enable
- Enable_all
- Drop
- Drop_all
- Show_filters
- Alter
- Alter_status

Data manipulation commands

These commands will work on the table related to data manipulations such as putting data into a table, retrieving data from a table and deleting schema, etc.

The commands come under these are

- Count
- Put
- Get
- Delete
- Delete all
- Truncate
- Scan